

Image-Based View Synthesis by Combining Trilinear Tensors and Learning Techniques

S. Avidan T. Evgeniou A. Shashua T. Poggio *

Hebrew University
Institute of Computer Science
Jerusalem 91904
Israel

Artificial Intelligence Laboratory
M.I.T
Cambridge, MA 02139
USA

Abstract

We present a new method for rendering novel images of flexible 3D objects from a small number of example images in correspondence. The strength of the method is the ability to synthesize images whose viewing position is significantly far away from the viewing cone of the example images (“view extrapolation”), yet without ever modeling the 3D structure of the scene. The method relies on synthesizing a chain of “trilinear tensors” that governs the warping function from the example images to the novel image, together with a multi-dimensional interpolation function that synthesizes the non-rigid motions of the viewed object from the virtual camera position. We show that two closely spaced example images alone are sufficient in practice to synthesize a significant viewing cone, thus demonstrating the ability of representing an object by a relatively small number of model images — for the purpose of cheap and fast viewers that can run on standard hardware.

1 Introduction

In this paper we develop a reprojection technique for synthesizing novel views of a 3D object, given a collection of 2D model images in correspondence. We also consider the synthesis of novel views from a single model image and develop a method to control non-rigid transfor-

mations of the viewed object, such as facial expressions in the case of faces.

The most significant aspect of our approach is the ability to synthesize images that are far away from the viewing positions of the sample model images without ever computing explicitly any 3D information about the scene. This property provides a multi-image representation of the 3D object using a minimal number of images. In our experiments, for example, two closely spaced frontal images of a face are sufficient for generating photo-realistic images from viewpoints within a 60 degrees cone of visual angle – further extrapolation is possible but the image quality degrades. The immediate application of our results is to provide a very fast 3D viewing system based on a small number of images that can run on standard hardware.

The notion of image-based rendering is gaining momentum in both the computer graphics and computer vision communities. The general idea is to achieve photo-realistic virtual images while avoiding the computational-intensive process of acquiring a 3D model followed by rendering. Instead, one seeks to use a number of model images of the 3D object or scene as a representation from which novel views can be synthesized *directly* by means of image warping.

The forerunner of this approach is to create a panoramic image of a scene (mosaic) from overlapping images taken from a fixed location while varying the orientation of the camera. The mosaic is mapped to a virtual cylinder that allows the user to look continuously at all directions but not to move. This is the basis for the QuickTimeVR system [7].

The fixed position constraint can be relaxed by computing the optical flow between the example images and using it to interpolate between the cylinders constructed at different locations (cf. [4, 5, 6]) (originally proposed

*S. Avidan and A. Shashua are with the Hebrew University. T. Evgeniou and T. Poggio are with M.I.T. This research is sponsored by a ONR grant under contract N00014-96-10342 and by a grant from the National Science Foundation under contract ASC-9217041. Additional support is provided by Eastman Kodak Company, Siemens Corporate Research, Inc., ATR and AT&T. A full version of this paper appeared as TR-1603, MIT AI Lab, 1997

for views, not mosaics, but the principle is the same). However, interpolation may produce physically-invalid images. Seitz and Dyer [20] proposed a physically-valid view interpolation method. The method involves recovering the epipolar geometry between the two acquired images and having interpolation done along the rectified epipolar lines.

Interpolation can also be performed directly on the plenoptic function [1] which represents the amount of light emitted at each point in space as a function of direction. Levoy et al. [15] and Gortler et al. [10] interpolate between a dense set of several thousands of example images to reconstruct a reduced plenoptic function (under an occlusion-free world assumption). They considerably increase the number of example images to avoid computing optical flow between the model images.

The major limitations of the aforementioned techniques is that a relatively large number of model images is required to represent an object. The alternative approach, along the lines of this paper, is to reduce the number of acquired (model) images by exploiting the 3D-from-2D geometry of the problem with the aid of corresponding points between the model images. Laveau and Faugeras [14] were the first to use the epipolar constraint for view synthesis, allowing them to extrapolate, as well as interpolate, between the example images. Epipolar constraints, however, are subject to singularities that arise under certain camera motions (like when the virtual camera center is collinear with the centers of the model cameras) and the relation between translational and rotational parameters of the virtual camera and the epipolar constraint is somewhat indirect and hence requires the specification of matching points [14]. The singular camera motions can be relaxed by using the depth map of the environment. McMillan and Bishop [17] use a full depth map (3D reconstruction of the camera motion and the environment) together with the epipolar constraint to provide a direct connection between the virtual camera motion and the re-projection engine. Depth maps are easily provided for synthetic environments, whereas for real scenes the process is fragile especially under small base-line situations that arise due to the requirement of dense correspondence between the model images/mosaics [11].

In this paper we propose a new view-synthesis method that makes use of the recent development of multi-linear matching constraints, known as trilinearities, that were first introduced in [21]. The trilinearities provide a general (not subject to singular camera configurations) warping function from model images to novel synthesized images governed directly by the camera parameters of the virtual camera. Therefore, we provide a true multi-image

system for view synthesis that does not require a companion depth map, nor the full reconstruction of camera parameters among the model cameras, yet is general and robust. The strength of our method is demonstrated by the ability to work with closely spaced acquired images yet synthesize high-quality views at a significant extrapolation from the viewing angles of the acquired images. Furthermore, our method can be generalized to work with a single acquired model image and to allow non-rigid transformations by integrating multi-linear constraints and multi-dimensional interpolation.

The main contributions of our work are:

1. The introduction of the trilinear tensor as the warping function.
2. The derivation of a tensorial operator which is the heart of the method. The operator generates a cascading set of tensors from two model views in correspondence and the parameters of the virtual camera motion. The tensorial operator does not require the estimation of the baseline between the acquired images (typically a fragile process) thereby enabling the model images to be closely spaced without hindering much the robustness of the synthesis process.
3. The combination of the tensor with a learning method for the generation of virtual views of an object given only a single model image.
4. The combination of rigid transformations using the tensor with non-rigid transformations achieved with multi-dimensional view interpolation.

On the experimental side, we have tested the proposed method on a variety of objects with a variety of cameras in real-world conditions where neither camera calibration (or even camera type) is available nor the lighting conditions are controlled. We demonstrated that correspondence is practical for closely spaced images, and that the synthesis method is sufficiently accurate and robust.

2 View Synthesis in Tensor Space

The view synthesis approach is based on the following paradigm. Three views satisfy certain matching constraints of a trilinear form, represented by a tensor. Thus, given two views in correspondence and a tensor, the corresponding third view can be generated uniquely by means of a warping function, as described below in more detail. We describe how to recover the tensor parameters and show a “driver” function that governs the change in tensor coefficients as a result of moving the virtual camera.

2.1 The Trilinear Warping Function

The trilinear tensor concatenates together the camera transformation matrices (camera locations) across three views, as follows. Let P be a point in 3D projective space projecting onto p, p', p'' in three views ψ, ψ', ψ'' respectively, represented by the two dimensional projective space. The relationship between the 3D and the 2D spaces is represented by the 3×4 matrices, $[I, 0]$, $[A, v']$ and $[B, v'']$, i.e.,

$$\begin{aligned} p &= [I, 0]P \\ p' &\cong [A, v']P \\ p'' &\cong [B, v'']P \end{aligned}$$

where A, B stand for the rotational component of camera motion (generally these are 2D homography matrices) and v', v'' stand for the translational component (generally these are the epipolar points).

We may adopt the convention that $p = (x, y, 1)^\top$, $p' = (x', y', 1)^\top$, $p'' = (x'', y'', 1)^\top$ and, thus, $P = (x, y, 1, \rho)$. The coordinates $(x, y), (x', y'), (x'', y'')$ are matching points across the three images.

The trilinear tensor is an array of 27 entries:

$$\alpha_i^{jk} = v'^j b_i^k - v''^k a_i^j, \quad i, j, k = 1, 2, 3 \quad (1)$$

where the covariant-contravariant indexing notation is assumed (see Appendix A). The tensor α_i^{jk} forms the set of coefficients of certain trilinear forms that vanish on any corresponding triplet p, p', p'' :

$$p^i s_j^\mu r_k^\rho \alpha_i^{jk} = 0 \quad (2)$$

where s_j^μ are any two lines (s_j^1 and s_j^2) intersecting at p' , and r_k^ρ are any two lines intersecting at p'' (see Fig. 1).

Since each of the free indices μ, ρ is in the range 1,2, we have 4 trilinear equations which are unique up to linear combinations. If we choose the canonical form where s and r represent vertical and horizontal lines, then the four trilinear forms, referred to as trilinearities, are expanded as follows:

$$\begin{aligned} x'' \alpha_i^{13} p^i - x'' x' \alpha_i^{33} p^i + x' \alpha_i^{31} p^i - \alpha_i^{11} p^i &= 0, \\ y'' \alpha_i^{13} p^i - y'' x' \alpha_i^{33} p^i + x' \alpha_i^{32} p^i - \alpha_i^{12} p^i &= 0, \\ x'' \alpha_i^{23} p^i - x'' y' \alpha_i^{33} p^i + y' \alpha_i^{31} p^i - \alpha_i^{21} p^i &= 0, \\ y'' \alpha_i^{23} p^i - y'' y' \alpha_i^{33} p^i + y' \alpha_i^{32} p^i - \alpha_i^{22} p^i &= 0. \end{aligned}$$

Since every corresponding triplet p, p', p'' contributes four linearly independent equations, then seven corresponding points across the three views uniquely determine (up to scale) the tensor α_i^{jk} . These constraints first




Figure 1: Each of the four trilinear equations describes a matching between a point p in the first view, some line s_j^μ passing through the matching point p' in the second view and some line r_k^ρ passing through the matching point p'' in the third view. In space, this constraint is an intersection between a ray and two planes.

became prominent in [21] and the underlying theory has been studied intensively in [24, 12, 9, 25, 13, 22].

One can readily see that given two views in full correspondence and the tensor (recovered using 7 matching points with a third view), the entire third view can be synthesized by means of forward warping. From each trilinearity we can simply extract either x'' or y'' , thus for every matching pair p, p' we can obtain p'' . We then copy at p'' the appropriate brightness value, for example the average of the pixel values at p and p' in the two model images. This process is referred to as “reprojection” in the literature. There are alternative ways of performing reprojection, but if we would like to do it without recovering first a 3D model of the scene, the trilinear tensor generally provides the best results since it is free from singular configurations (see [2, 21, 23]). In image-based rendering we would like to obtain the tensor via user specification of the location of a virtual camera, rather than by the specification of (at least) seven matching points. This is described next.

2.2 The basic Tensorial Operator

The basic tensorial operator describes how to modify (transform) a tensor so as to represent a new configuration of three cameras. We are particularly interested in the case where only one camera has changed its position and orientation. Thus, by repeated application of the operator on a seed tensor with a sequence of desired virtual camera positions (translation and orientation) we obtain a chain of warping functions (tensors) from the set of

acquired images (from which the seed tensor was computed) to create the desired virtual views.

Consider the tensor α_i^{jk} of the views $\langle 1, 2, 3 \rangle$ (in that order), and assume the user wishes to apply an incremental change of position of the third image, i.e., rotate the third camera position by the 3×3 coordinate matrix C , and translate it by the 3×1 translation vector t — this motion would result to a novel view, call it view 4. Then the tensor γ_i^{jk} of the views $\langle 1, 2, 4 \rangle$ is given by:

$$\gamma_i^{jk} = v'^j (c_l^k b_i^l) - (v''^l c_l^k + t^k) a_i^j = c_l^k \alpha_i^{jl} - t^k a_i^j. \quad (3)$$

This is so because we use Eq. 1 where we replace the motion parameters v''^k, b_i^k from the first image to the third image, with $(v''^l c_l^k + t^k), (c_l^k b_i^l)$ which depend on the motion parameters from the first image to the novel one. The matrix a_i^j representing the rotational component of camera motion between the two model views 1,2 can be represented in closed form as a function of the tensor α_i^{jk} as described in [19]:

$$\begin{aligned} \Omega_X &= \det \begin{pmatrix} \alpha_2^{j3} & \alpha_3^{j3} \\ \alpha_2^{j3} + \alpha_3^{j2} & \alpha_3^{j3} - \alpha_2^{j2} \end{pmatrix} / K \\ \Omega_Y &= \det \begin{pmatrix} -\alpha_1^{j3} & \alpha_2^{j3} \\ \alpha_2^{j3} + \alpha_3^{j2} & \alpha_3^{j3} - \alpha_2^{j2} \end{pmatrix} / K \\ \Omega_Z &= \det \begin{pmatrix} \alpha_1^{j2} & \alpha_2^{j3} \\ \alpha_2^{j3} + \alpha_3^{j2} & \alpha_3^{j3} - \alpha_2^{j2} \end{pmatrix} / K \\ K &= \det \begin{pmatrix} \alpha_2^{j2} & \alpha_2^{j3} \\ \alpha_2^{j3} + \alpha_3^{j2} & \alpha_3^{j3} - \alpha_2^{j2} \end{pmatrix} \end{aligned} \quad (4)$$

where $\Omega_X, \Omega_Y, \Omega_Z$ are rotation angles and α_2^{j2} stands for $(\alpha_2^{12}, \alpha_2^{22}, \alpha_2^{32})$, etc.

To summarize, Eq. 3 is a general formula for transforming the tensor based on an incremental camera motion of a fixed (third) camera. Therefore, starting from a “seed” tensor and a sequence of desired camera motions, the set of corresponding tensors can be generated and used to warp the acquired images onto the novel views. We next consider how to obtain the seed tensor that starts the process.

2.3 The Seed Tensor of Two Views

Given two acquired images we can construct a special tensor composed of the elements of the fundamental matrix [8] that can serve as a seed tensor that starts the chain of tensors, as follows. Consider a configuration of three

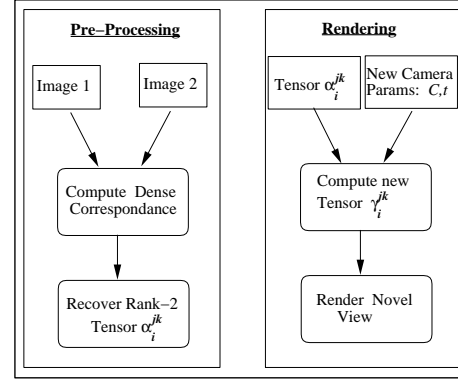


Figure 2: View synthesis is divided into two parts. The pre-processing stage, done only once and the actual rendering done for every image.

views in which views 2,3 coincide, i.e., Eq. 1 becomes:

$$\gamma_i^{jk} = v'^j a_i^k - v'^k a_i^j \quad (5)$$

where γ_i^{jk} is the tensor of the image triplet $\langle 1, 2, 2 \rangle$. It can be readily verified that the elements of γ_i^{jk} are composed of the fundamental matrix $f_{ij} = \epsilon_{ikl} v'^k a_l^j, -f_{ij}$, and the remaining (nine) elements vanish. It will not be shown here, but the rank of γ_i^{jk} is 2 whereas the rank of the tensor of three distinct views is 4 — but otherwise all other properties remain and, in particular, γ_i^{jk} can serve as the first tensor that starts the synthesis process described above.

2.4 The View Synthesis Loop

The method is divided into two stages — a preprocessing stage, done only once, and the actual rendering done for every new frame (see Fig. 2). In the preprocessing stage compute dense correspondence between the two model images, recover the fundamental matrix and construct the rank-2 tensor (eq 5) from it. In rendering time, accept virtual camera parameters from the user, generate the new tensor (eq 3) and synthesize the novel view from the two model images and the computed tensor.

2.5 Experiments

We conducted our experiments on pairs of model images, taken with uncalibrated cameras. Rendering time is approximately 5 seconds on an SGI Indy machine, without any optimization. Due to space limitations, we show here only the “Shannon” example (Figure 3), where a pair of images, 620×764 pixels each, were taken and extrapolated up to 90 degrees with graceful degradation in image quality.




Figure 3: The “Shannon” example of view synthesis using our reprojection method. The only real images are the two framed (dimmed) ones on the top row. Our method allows both interpolation (center image) and extrapolation, as can be seen in the lower two rows.

3 Synthesis from a single model view

In the previous section, we have discussed how to synthesize new images for different viewpoints, given two examples images. Suppose now that only one image of a specific object, say \mathbf{p}_{nov} , is available. Does our reprojection method break down? This section sketches a solution to this question. The idea is to create an additional – virtual – example image of object \mathbf{p}_{nov} from just one real view of it. Once one obtains a second image of the object from a different viewpoint, one can use the reprojection algorithm of section 2.4 to generate subsequent virtual images. To accomplish this task without using a parametric 3D model, one may start from a collection of example views of another similar object \mathbf{p} which plays the role of a prototype for representing generic transformations of the object class that the two objects belong to. Faces form such a class of objects. In general, we want to generate from one 2D view \mathbf{Img}_{nov} of a 3D object \mathbf{p}_{nov} other views, exploiting knowledge of views of other objects of the same class. This idea of generating “virtual” views of an object by using class-specific knowledge has been discussed before (see references in [5]). Suppose that we have two views \mathbf{Img}_{ref} and \mathbf{Img}_p of the prototype. We take \mathbf{Img}_{ref} to appear in the same pose as \mathbf{Img}_{nov} . \mathbf{Img}_p is a slightly transformed (i.e., rotated) view of \mathbf{Img}_{ref} (see diagram in Fig. 4). We can then compute the optical flow \mathbf{S}_p between these two views. Moreover, since the prototype object \mathbf{p} is assumed to be

“similar” to object \mathbf{p}_{nov} , we assume that we can find good correspondence \mathbf{S}_{nov} between \mathbf{Img}_{ref} and \mathbf{Img}_{nov} . We subsequently generate the optical flow \mathbf{S}_{p+nov} between the view \mathbf{Img}_{ref} and a new view of the object \mathbf{p}_{nov} by the vector addition:

$$\mathbf{S}_{p+nov} = \mathbf{S}_p + \mathbf{S}_{nov}. \quad (6)$$

A new view, \mathbf{Img}_{p+nov} , of object \mathbf{p}_{nov} is then rendered by texture mapping from the single available view \mathbf{Img}_{nov} after forward warping from \mathbf{Img}_{ref} using optical flow \mathbf{S}_{p+nov} . In a sense, we “map” the learned transformation (optical flow \mathbf{S}_p) from \mathbf{Img}_{ref} to \mathbf{Img}_{nov} using flow \mathbf{S}_{nov} . We now have two images of object \mathbf{p}_{nov} that our reprojection technique can use to simulate a virtual camera and generate new images and image sequences.

We demonstrate this technique in Fig. 4 using as an example a self-portrait of Van Gogh. A slight rotation is learned from another similar prototypical “object”, in this case another face, to generate a first virtual image of Van Gogh. Then the reprojection method of section 2.4 is used to generate subsequent views.

4 Incorporating non-rigid Transformations

So far we have described a technique that allows the user to generate new images by controlling the rigid degrees of freedom that correspond to motion of the camera. From two or more images in correspondence it may

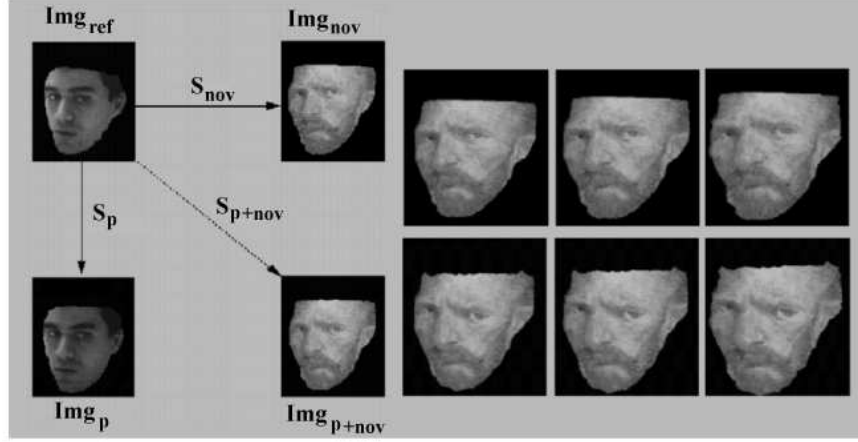


Figure 4: We can simulate a virtual camera looking at an object from arbitrary viewpoints given only one image. First we synthesize a single virtual view of the object by learning an appropriate transformation, in this case a small rotation, from another similar prototypical object. The method is shown on the left diagram. In this example we create a rotated self-portrait of Van Gogh (Img_{p+nov}) using the original self-portrait (Img_{nov}) and information from a “similar” prototypical face (Img_{ref} and Img_p). S_{nov} is the optical flow between the first image of the prototypical face and the original painting. S_p is the optical flow between the two images of the prototypical face. The new Van Gogh image is then created by texture mapping from the single available painting after forward warping from Img_{ref} using $S_{p+nov} = S_p + S_{nov}$. From these two views of Van Gogh we then generate other views from arbitrary viewpoints, as shown on the right, using our reprojection method.

also be possible to generate new images of non-rigid 3D objects as a function of input parameters that correspond to non-rigid transformations such as a change in expression of a face. The underlying operation is multidimensional interpolation, a simple extension of traditional image morphing. We outline the technique and illustrate how it can be integrated with the algebraic method described so far.

Let us assume that n images are available and that pixel-wise correspondence can be computed with an optical flow algorithm between one of them, chosen as the “zero” reference image, and each of the $n - 1$ others. As we saw earlier, correspondence associates to each image i the optical flow, that we note as S_i , of the position of each pixel relative to the reference image. We can also associate to each image a vector of color values, the “texture” vector, that we note as T_i . The texture vector T_i is simply the image i warped to the shape of the reference image by the optical flow S_i . Let us also assume that the user defines the values r_i of the non-rigid parameters of interest to be associated with each one of these “example” images.

A multidimensional interpolation technique such as Radial Basis Functions or splines is then used to interpolate the n example pairs $(r_i, (S_i, T_i))$ (see for instance

[5]). The mapping from the input space of non-rigid parameters to the output space of images, expressed in terms of textures and flows, is provided by the following interpolation scheme which can be regarded as a learning network ([4, 5])

$$\begin{aligned} S(r) &= \sum_{i=1}^n c_i G(r - r_i), \\ T(r) &= \sum_{i=1}^n a_i G(r - r_i), \end{aligned} \quad (7)$$

where the c_i and a_i are vectors of coefficients, and G is a basis function, which may be a radial basis function, like the Gaussian or a spline, like a tensor product spline. The network coefficients c_i and a_i are found by solving the linear system of equations (7) over the training data ([4, 5]).

Given a new vector r of non-rigid parameters, the network of Eq. 7 synthesizes a new (S, T) using the learned coefficients c_i and a_i , which is then rendered in a new image by warping T_i according to the warping field S_i , effectively performing multidimensional morphing.

This simple technique can be used to control several non-rigid degrees of freedom such as facial expressions, as shown by [4, 5]. It can be combined directly with the algebraic technique described earlier to control the position of the virtual camera. That is, we first perform



Figure 5: We control non-rigid degrees of freedom by combining interpolation techniques with our reprojection method. Using the four example images on the left (framed) we generate new images as shown on the right. We can synthesize intermediate expressions from any viewpoint.

a multi-dimensional interpolation (non-rigid transformation) followed by the image extrapolation (rigid transformation) described in section 2.4.

Fig. 5 shows the simple case of one non-rigid degree of freedom. Given four images, corresponding to two values of the non-rigid parameter and two viewpoints, a virtual image for the desired intermediate expression is obtained by interpolation for each of the viewpoints. Then our reprojection technique generates views of the intermediate expression from the desired new viewpoint.

5 Conclusions

The method we describe in this paper can render novel images of flexible 3D objects from a small number of example images without the need of an explicit 3D model. Its main strength is the ability to synthesize images whose viewing position is significantly far away from the viewing cone of the example images.

Clearly the key step in this class of techniques is the computation of pixel-wise correspondence between the example images. We addressed this problem by using an optical flow algorithm from the computer vision literature that estimates dense sets of pixel-level correspondences. It is well known that correspondence is a very difficult problem which can be solved only for images that are similar enough and do not suffer from significant self-occlusions. 3D model-based approaches, however, suffer from even worse correspondence problems, if the 3D models are themselves estimated from a set of images. As the experimental results indicate, our technique has a major advantage relative to others, at least for the rigid degrees of freedom of the camera, since it relies on pairs of images with a *small* baseline, which helps the critical correspondence stage. Methods that estimate 3D struc-

ture are very noisy with small baselines. Although our technique implicitly estimates the 3D structure, not doing so in an explicit way means avoiding noisy steps and therefore generating less noisy virtual images. Moreover, morphing techniques, such as [20], require large baselines since they cannot perform extrapolation.

References

- [1] Adelson, E.H. and J.R. Bergen. The Plenoptic Function and the Elements of Early Vision. In *Computational Models of Visual Processing*, Chapter 1, Edited by Michael Landy and J. Anthony Movshon. The MIT Press, Cambridge, Mass. 1991.
- [2] E.B. Barrett, P.M. Payton, and G. Gheen. Robust algebraic invariant methods with applications in geometry and imaging. In *Proceedings of the SPIE on Remote Sensing*, San Diego, CA, July 1995.
- [3] J.R. Bergen, P. Anandan, K.J. Hanna, and R. Hingorani. Hierarchical model-based motion estimation. In *Proceedings of the European Conference on Computer Vision*, Santa Margherita Ligure, Italy, June 1992.
- [4] D. Beymer, A. Shashua, and T. Poggio. Example based image analysis and synthesis. A.I. Memo No. 1431, Artificial Intelligence Laboratory, Massachusetts Institute of Technology, October 1993.
- [5] D. Beymer, and T. Poggio. Image Representations for Visual Learning. *Science*, 272, pages 1905–1909, 1996.
- [6] S.E. Chen and L. Williams. View interpolation for image synthesis. In *SIGGRAPH*, pages 279–288, Anaheim, CA, August 1993.
- [7] Shenchang Eric Chen. QuickTimeVR - an image-based approach to virtual environment navigation. In *SIGGRAPH*, pages 29–38, 1995.
- [8] O.D. Faugeras. What can be seen in three dimensions with an uncalibrated stereo rig? In *Proceedings of the European Conference on Computer Vision*, pages 563–578, Santa Margherita Ligure, Italy, June 1992.
- [9] O.D. Faugeras and B. Mourrain. On the geometry and algebra of the point and line correspondences between N images. In *Proceedings of the International Conference on Computer Vision*, Cambridge, MA, June 1995.

- [10] Gortler, S.J., Grzeszczuk, R., Szeliski, R. and Cohen, M. The Lumigraph In *SIGGRAPH*, pages 43–54, 1996.
- [11] W.E.L. Grimson. Why stereo vision is not always about 3D reconstruction. A.I. Memo No. 1435, Artificial Intelligence Laboratory, Massachusetts Institute of Technology, July 1993.
- [12] R. Hartley. A linear method for reconstruction from lines and points. In *Proceedings of the International Conference on Computer Vision*, pages 882–887, Cambridge, MA, June 1995.
- [13] A. Heyden. Reconstruction from image sequences by means of relative depths. In *Proceedings of the International Conference on Computer Vision*, pages 1058–1063, Cambridge, MA, June 1995.
- [14] S. Laveau and O.D. Faugeras. 3-d scene representation as a collection of images. In *Proceedings of the International Conference on Pattern Recognition*, pages 689–691, Jerusalem, Israel, October 1994.
- [15] M. Levoy and P. Hanrahan. Light field rendering. In *SIGGRAPH*, pages 31–42, August 1996.
- [16] B.D. Lucas and T. Kanade. An iterative image registration technique with an application to stereo vision. In *Proceedings IJCAI*, pages 674–679, Vancouver, Canada, 1981.
- [17] Leonard McMillan and Gary Bishop. Plenoptic modeling: An image-based rendering system. In *SIGGRAPH*, pages 39–46, 1995.
- [18] Torr P.H.S., Zisserman A., and Murray D. Motion clustering using the trilinear constraint over three views. In *Workshop on Geometrical Modeling and Invariants for Computer Vision*. Xidian University Press., 1995.
- [19] B. Rousso, S. Avidan, A. Shashua, and S. Peleg. Robust recovery of camera rotation from three frames. In *Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition*, San-Francisco, CA, June 1996.
- [20] Steven M. Seitz, Charles R. Dyers. View morphing. In *SIGGRAPH*, pages 21–30, 1996.
- [21] A. Shashua. Algebraic functions for recognition. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 17(8):779–789, 1995.
- [22] A. Shashua and S. Avidan. The rank4 constraint in multiple view geometry. In *Proceedings of the European Conference on Computer Vision*, Cambridge, UK, April 1996.
- [23] A. Shashua and S.J. Maybank. Degenerate n point configurations of three views: Do critical surfaces exist? Technical Report TR 96-19, Hebrew University of Jerusalem, November 1996.
- [24] A. Shashua and M. Werman. On the trilinear tensor of three perspective views and its underlying geometry. In *Proceedings of the International Conference on Computer Vision*, June 1995.
- [25] B. Triggs. Matching constraints and the joint image. In *Proceedings of the International Conference on Computer Vision*, pages 338–343, Cambridge, MA, June 1995.

A On Tensorial Notations

We use the covariant-contravariant summation convention: a point is an object whose coordinates are specified with superscripts, i.e., $p^i = (p^1, p^2, \dots)$. These are called contravariant vectors. An element in the dual space (representing hyper-planes — lines in \mathcal{P}^2), is called a covariant vector and is represented by subscripts, i.e., $s_j = (s_1, s_2, \dots)$. Indices repeated in covariant and contravariant forms are summed over, i.e., $p^i s_i = p^1 s_1 + p^2 s_2 + \dots + p^n s_n$. This is known as a contraction. For example, if p is a point incident to a line s in \mathcal{P}^2 , then $p^i s_i = 0$. Vectors are also called 1-valence tensors. 2-valence tensors (matrices) have two indices and the transformation they represent depends on the covariant-contravariant positioning of the indices. For example, a_i^j is a mapping from points to points, and hyper-planes to hyper-planes, because $a_i^j p^i = q^j$ and $a_i^j s_j = r_i$ (in matrix form: $Ap = q$ and $A^T s = r$). When viewed as a matrix the row and column positions are determined accordingly: in a_i^j and a_j^i the index i runs over the columns and j runs over the rows, thus $b_j^k a_i^j = c_i^k$ is $BA = C$ in matrix form. An outer-product of two 1-valence tensors (vectors), $a_i b^j$, is a 2-valence tensor c_i^j whose i, j entries are $a_i b^j$ — note that in matrix form $C = ba^T$. The tensor of vector products is denoted by ϵ_{ijk} (indices range 1-3) operates on two contravariant vectors of the 2D projective plane and produces a covariant vector in the dual space (a line): $\epsilon_{ijk} p^i q^j = s_k$, which in vector form is $s = p \times q$, i.e., s is the vector product of the points p and q .