

# CAP 5415 – Computer Vision

Marshall Tappen  
Fall 2010

Lecture 1

# Welcome!

- About Me

- 
- Interested in Machine Vision and Machine Learning
- Happy to chat with you at almost any time
  - May want to e-mail me first
- Office Hours:
  - Tuesday-Thursday before class

# Grading

- Problem Sets – 50%
- 3 Solo Problem Sets – 50%
  - You may not collaborate on these

# Doing the problems

- Finishing the problem sets will require access to an interpreted environment
  - MATLAB
  - Octave
  - Numerical Python
- **NO COMPILED LANGUAGES!!!!**
  - No C/C++
  - No Java
  - No x86 Assembler
- My Compiled Languages Rant

# Environments

- MATLAB
  - Pro: Well-established package. You can find many tutorials on the net.
  - Con: Not free. If your lab does not already have it, talk to me about getting access.
- Octave
  - Free MATLAB look-alike
  - Pro: Should be able to handle anything you will do in this class
  - Con: “Should be”. I'm not sure about support in Windows

# Environments

- Numerical Python
  - All the capabilities of MATLAB
  - Free!
  - Real programming language
  - Used for lots of stuff besides numerical computing
  - Cons: Documentation is a bit sparse and can be outdated
    - I can get you started – I am working on a tutorial

# Math

- We will use it
- We will be talking about mathematical models of images and image formation
- This class is not about proving theorems
- My goal is to have you build intuitions about the models
- Try and visualize the computation that each equation is expressing
- Basic Calculus and Basic Linear Algebra should be sufficient

# Course Text

- We will use Szeliski Book – Free this year!
- Not required – more of a reference




# Course Structure

- This year, we will be covering pattern recognition more deeply than in previous years
- Machine learning is critical to modern computer vision
- You need to understand it well
- Important Foundational Topics:
  - Image Processing
  - Optimization
  - Machine Learning
  - Geometry


# Image Processing

- For now, we won't worry about the physical aspects of getting images
- View image as an array of continuous values




# Simple Modification


- What if we wanted to blur this image?
- We could take a local average
  - Replace each pixel with the mean of an  $N \times N$  pixel neighborhood surrounding that pixel.



# 3x3 Neighborhood




Original




Averaged

# 5x5 Neighborhood




Original




Averaged

# 7x7 Neighborhood





Original



Averaged

# Let's represent this more generally



0	3	0	0
0	6	1	16
0	0	2	46
0	0	2	43

# Let's represent this more generally

0	3	0	0
0	6	1	16
0	0	2	46
0	0	2	43

$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$
$1/9$	$1/9$	$1/9$




# Let's represent this more generally

0	3	0	0
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
0	6	1	16
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
0	0	2	46
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$
0	0	2	43
$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$	$\frac{1}{9}$

?	?	?	?
?	1.33	?	?
?	?	?	?
?	?	?	?

- Multiply corresponding numbers and add

# Let's represent this more generally



- Multiply corresponding numbers and add
- Template moves across the image
- Think of it as a sliding window

# This is called convolution

- Mathematically expressed as

$$R(i, j) = \sum_{m=-N}^N \sum_{n=-N}^N I(m, n) K(i - m, j - n)$$

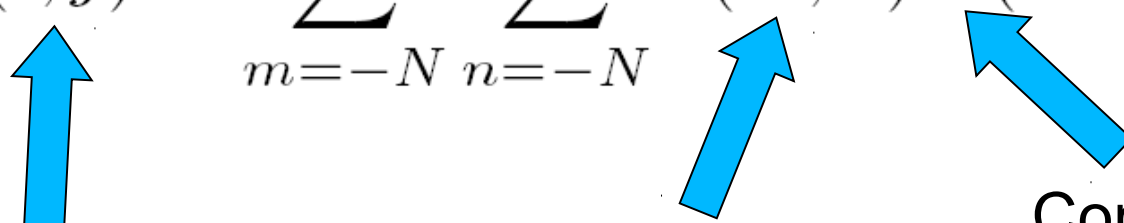
Resulting Image

Input Image

Convolution Kernel

# Take out a piece of paper

- Let's say  $i=10$  and  $j=10$
- Which location in  $K$  is multiplied by  $I(5,5)$ ?
- $I(5,4)$

$$R(i, j) = \sum_{m=-N}^N \sum_{n=-N}^N I(m, n) K(i - m, j - n)$$


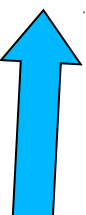
The diagram illustrates the convolution equation with three blue arrows pointing from labels below to terms in the equation above:

- An arrow points from "Resulting Image" to  $R(i, j)$ .
- An arrow points from "Input Image" to  $I(m, n)$ .
- An arrow points from "Convolution Kernel" to  $K(i - m, j - n)$ .


# Notation

- Also denoted as
- $R = I * K$
- We “convolve”  $I$  with  $K$ 
  - Not convolute!

$$R(i, j) = \sum_{m=-N}^N \sum_{n=-N}^N I(m, n) K(i - m, j - n)$$



Resulting Image



Input Image



Convolution Kernel

# Sliding Template View

- Take the template  $K$


1	2	3
4	5	6
7	8	9

- Flip it

9	8	7
6	5	4
3	2	1


- Slide across image

# Predict the Image




0	0	0
0	1	0
0	0	0

# Predict the Image




0	0	0
0	1	0
0	0	0






# Predict the Image




0	0	0
0	0	1
0	0	0


# Predict the Image




0	0	0
0	0	1
0	0	0




# Predict the Image



# Predict the Image




# Predict the Image




1	0	0
0	0	0
0	0	0


# Predict the Image




1	0	0
0	0	0
0	0	0



# Predict the Image




# Predict the Image






# Predict the Image




0	1	0
1	2	1
0	1	0

# Predict the Image



0	1	0
1	2	1
0	1	0



# Predict the kernel

- What if I wanted to compute

$$R(i,j) = I(i+1,j) - I(i,j)$$

at every pixel?

- What would the kernel be?
- This is one discrete approximation to the derivative

# What's the problem with this derivative?

$[1 \ -1 \ 0]$

– Where's the center of the derivative?

- An alternative

–  $[-1 \ 0 \ 1]$

# Your Convolution filter toolbox


- In my experience, 90% of the filtering that you will do will be either
  - Smoothing (or Blurring)
  - High-Pass Filtering (I'll explain this later)
- Most common filters:
  - Smoothing: Gaussian
  - High Pass Filtering: Derivative of Gaussian

# Gaussian Filter

- Let's assume that a  $(2k+1) \times (2k+1)$  filter is parameterized from  $-k$  to  $+k$
- The Gaussian filter has the form

$$K(i, j) = \frac{1}{Z} \exp \left( -\frac{i^2 + j^2}{2\sigma^2} \right)$$

- And looks like




# Derivative of Gaussian Filter

- Take the derivative of the filter with respect to  $i$ :

$$\frac{\partial K(i, j)}{\partial i} = \frac{-i}{\sigma^2 Z} \exp \left( -\frac{i^2 + j^2}{2\sigma^2} \right)$$


- Filter looks like:




- Basically blur then take the derivative

# Effect of Changing $\sigma$

- With  $\sigma$  set to 1
- With  $\sigma$  set to 3



Input





# Practical Aspects of Computing Convolutions


- Let's blur this flat, gray image:




- What should it look like?

# Practical Aspects of Computing Convolutions

- Depending on how you do the convolution in MATLAB, you could end up with 3 different images



266x266 Image



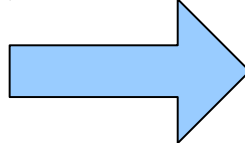
256x256 Image



246x246 Image

# Border Handling

- Lets go back to the sliding template view
- What if I wanted to compute an average right here?




0	3	0	0
0	6	1	16
0	0	2	46
0	0	2	43

# Border Handling

1/9	1/9	1/9		
1/9	1/9	1/9	0	0
1/9	1/9	1/9	1	16
	0	0	2	46
	0	0	2	43

# Practical Aspects of Computing Convolutions

- Filled in borders with zeros, computed everywhere the kernel touches
- Called “full” in MATLAB



266x266 Image

# Border Handling

1/9	1/9	1/9			
1/9	1/9	1/9			
1/9	1/9	1/9	0	3	0
			0	6	1
			0	0	2
			0	0	2
					16
					46
					43

# Practical Aspects of Computing Convolutions

- Fill in border with zeros, only compute at “original pixels”
- Called “same” in MATLAB



256x256 Image

# Border Handling

1/9	1/9	1/9			
1/9	0	3	0		0
1/9	0	6	1		16
	0	0	2		46
	0	0	2		43



# Practical Aspects of Computing Convolutions

- Only compute at places where the kernel fits in the image



246x246 Image

# There are other options



- The first two methods that I described fill missing values in by substituting zero
- Can fill in values with different methods
  - Reflect image along border
  - Pull values from other side
- Not supported in MATLAB's convolution
  - Eero Simoncelli has a package that supports that kind of convolution

# Going Non-Linear

- Convolution is a linear operation
  - What does that mean?
- A simple *non-linear* operation is the median filter
  - Will explore that filter in the first problem set.

# Practical Use of These Properties – Image Sharpening

- Take this image and blur it



# Basic Convolution Properties

$$f(x) * ((g(x) * h(x))) = (f(x) * g(x)) * h(x)$$

$$(\alpha f(x)) * g(x) = \alpha (f(x) * g(x))$$

$$f(x) * g(x) = g(x) * f(x)$$


$$f(x) * (g(x) + h(x)) = f(x) * g(x) + f(x) * h(x)$$

- Can derive all of these with the definition of convolution
- Comes from linearity of convolution

$$R(i, j) = \sum_{m=-N}^N \sum_{n=-N}^N I(m, n) K(i - m, j - n)$$

# Practical Use of These Properties – Image Sharpening


- What do we get if we subtract the two?




- This is the leftover “sharp-stuff”

# Let's make the image sharper


- We know



+




=




Let's boost the sharp stuff a little

2x




+







Let's boost the sharp stuff a little



# Side-by Side



# Now look at the computation


- Operations
  - 1 convolution
  - 1 subtraction over the whole image
- As an equation:

$$\mathcal{I} * f + 2 (\mathcal{I} - \mathcal{I} * f)$$


Rewrite this

$$\mathcal{I} * f + 2(\mathcal{I} - \mathcal{I} * f)$$

$$\mathcal{I} * f + 2(\mathcal{I} * \delta - \mathcal{I} * f)$$



This is an identity filter or unit impulse



# Basic Convolution Properties

$$f(x) * ((g(x) * h(x))) = (f(x) * g(x)) * h(x)$$

$$(\alpha f(x)) * g(x) = \alpha (f(x) * g(x))$$

$$f(x) * g(x) = g(x) * f(x)$$

$$f(x) * (g(x) + h(x)) = f(x) * g(x) + f(x) * h(x)$$

- Can derive all of these with the definition of convolution
- Comes from linearity of convolution

$$R(i, j) = \sum_{m=-N}^N \sum_{n=-N}^N I(m, n) K(i - m, j - n)$$

Rewrite this

$$\mathcal{I} * f + 2(\mathcal{I} - \mathcal{I} * f)$$

$$\mathcal{I} * f + 2(\mathcal{I} * \delta - \mathcal{I} * f)$$

$$\mathcal{I} * (f + 2\delta - 2f)$$

Now look at the computation

$$\mathcal{I} * (f + 2\delta - 2f)$$

- Can pre-compute new filter
- Operations
  - 1 convolution